

DATA-260 Homework 1 Report

Name: Rohini Ramasheshu

Configuration Values


- SID4: 2604
- PORT_BASE: 8804
- PREFIX: s2604
- SEED: 2604
- VERIFY_SEED: 262604
- DOMAIN_ID: 4 (Open-source package vulnerabilities)

Hardware & Model

- Hardware: MacBook Air M2, 8GB RAM
- Local model used: qwen2.5:3b (substituted for qwen3:8b due to 8GB RAM constraint)
- Tagged commit hash: [INSERT AFTER TAGGING - see final step]

Part 1 — HTML/JavaScript Form

Built a vulnerability report form with required fields (package name, affected version, email, description, severity dropdown, terms checkbox) and JavaScript validation using arrow functions, JSON conversion, destructuring, spread operator, and closures for submission tracking.



Open-Source Package Vulnerability Report

Package Name AffectedVersion EmailId Description Severity ☒ I agree to the terms and conditions.

Deployment — Docker

Built and ran the app in a Docker container locally on port 8804.

A new version of Docker Desktop is available. Your current version is 146 days old.

Gordon

BETA

Containers

Images

Volumes

Kubernetes

Builds

Models

MCP Toolkit

BETA

Docker Hub

Docker Scout

Extensions

Containers

[Give feedback](#)

Container CPU usage

9.20% / 800% (8 CPUs available)

Container memory

2.1GB / 3.74GB

Search

Only show running containers


		Name	Container ID	Image	Port(s)
☐	☐	brave_bell	2120a793f56e	vuln-report	8804:8804
☐	☒	thirsty_hypatia	a413f9006e14	vuln-report	8804:8804 ↗
☐	>	mape-k-pipeline	-	-	-
☐	>	nosql	-	-	-

Deployment — AWS ECS

Deployed the Docker image via AWS ECR + ECS Fargate, exposing it on a public IP. Encountered and resolved a CPU architecture mismatch (image built for arm64 on Apple Silicon, but Fargate required amd64) by rebuilding with `--platform linux/amd64`.

file:///Users/rohiniramasheshu/data260-2604/reports/hw01/report.html

2/5



Open-Source Package Vulnerability Report

Package Name	e.g. lodash, express, re	AffectedVersion	e.g.4.2.5	EmailId	abc@gmail.com	Description	Describe the vulnerability in
Submit Vulnerability Report							

App running on AWS public IP

→ ↺ 🌐

us-east-2.console.aws.amazon.com/ecs/v2/clusters/data260-2604-cluster/services?region=us-east-2

Continue where you left off: Chrome can restore your tabs every time you restart. To turn this on, go to [chrome://flags/#enable-chrome-tabs-recovery](#)

Search

[Option+S]

Amazon Elastic Container Service > Clusters > data260-2604-cluster > Services

Amazon Elastic Container Service

favorites and recents

Express Mode

Clusters

Namespaces

Task definitions

Daemon task definitions

Account settings

WS Resource Explorer

WS Batch

Amazon ECR

Repositories

Online learning workshop

Documentation

Discover products

Subscriptions

Tell us what you think

Service updated: data260-2604-cluster:vuln-report-task-service-diayba1e

data260-2604-cluster

August 30, 2026, 11:55 PM

Cluster overview

Cluster ARN
arn:aws:ecs:us-east-2:491740041541:cluster/...

Status
Active

CloudWatch monitoring
Default

Services
1 active | 0 draining

Tasks
1 running | 0 pending

Registered container instances
-

Services

Tasks

Daemons

Infrastructure

Metrics

Scheduled tasks

Configuration - updated

Events

Services (1) Info

August 30, 2026, 19:46 (UTC-7:00)

Last updated

Filter services by value

Filter launch type
Any launch type

Filter scheduling strategy
Any scheduling strategy

☐	Service name	ARN	Status	Schedu...	Laun...	Task de...
☐	vuln-report-task-service-diayba1e	arn:aws:ec...	Active	Replica	-	vuln-repo...

Service and cluster were deleted immediately after confirming successful deployment to avoid ongoing AWS charges.

Part 2 — Agentic AI

Built a Planner → Reviewer → Finalizer pipeline using Ollama + langchain-ollama.

- Q1 (final tags): ["lodash", "security", "prototype-modification"]
- Q2 (final summary): "Prototype pollution vulnerability in lodash can lead to security issues in applications that handle user input." (19 words)
- Q3 (did Reviewer change anything?): Yes — it replaced the tag "vulnerability" with "security" for more standardized terminology, while keeping the summary otherwise consistent.

Explanation of each step: - **Planner**: given a title and content, drafts initial tags and a summary using a structured JSON-only prompt. - **Reviewer**: receives the Planner's draft and checks/improves it, returning JSON in the same format. - **Finalizer**: parses the Reviewer's JSON output, handling parse failures gracefully, and produces the final clean output.

Part 3 — Non-Determinism Experiment

See METRICS.md for full results table and raw/nondeterminism_results.csv for all 40 runs.

Summary: At temperature 0.7, 20 runs produced 10 distinct tag sets — meaningful run-to-run variation. At temperature 0.0, all 20 runs produced identical results (1 distinct tag set), confirming near-deterministic behavior at low temperature.

Part 4 — Model Client & Token Accounting

Built a reusable model_client.py adapter with a complete(messages, tools=None) interface, token counting per turn, and a /stats method reporting cumulative totals without altering history.

Ran a 5-turn conversation; /stats after turn 3 and turn 5 confirmed correct accumulation (turn 5: 5 turns, 194 total input tokens, 1848 total output tokens, 10662 history characters).

Observation: Turn 5 ("summarize what we discussed") produced a generic answer with no reference to earlier topics, revealing that our complete() method sends only the current message, not accumulated history — directly illustrating why real conversational memory requires explicitly resending context.

Theory answers: - *Why is prior context resent every turn?* LLMs are stateless between calls; without resending history, the model has no memory of prior turns, as demonstrated above. - *System prompt vs. user message?* A system prompt sets persistent behavior/instructions for the whole conversation; a user message is one turn of actual conversation content. - *Why do input tokens grow over a conversation?* If full history were resent each turn, each new turn's input would include all prior turns, growing linearly with conversation length. - *What eventually limits that growth?* The model's fixed context window — once history + new input exceeds it, older content must be truncated or summarized.

AI Use

See AI_USE.md for full disclosure.